

fp-tools: A Reproducible Platform for ATAC-seq Footprinting and Regulatory Motif Analysis

Yaoxiang Li ¹ and Chunling Yi ^{1,*}

¹ Department of Oncology, Lombardi Comprehensive Cancer Center, Georgetown University Medical Center, Georgetown University, Washington, DC, USA; yl814@georgetown.edu (Y.L.); cy232@georgetown.edu (C.Y.)

* Correspondence: cy232@georgetown.edu (C.Y.)

Abstract

Transcription-factor (TF) footprinting from ATAC-seq provides a base-resolution view of regulatory factor occupancy, yet complete analyses often require multiple tools with different interfaces, signal conventions, and reproducibility practices. We present *fp-tools*, a Python package that integrates Tn5 bias correction, footprint calling, motif matching, differential footprint analysis, and aggregate visualization within a consistent command-line workflow. The same framework supports de novo motif-discovery preparation, pseudobulk aggregation from single-cell ATAC fragments, and replicate-aware reporting for condition comparisons. Using public bulk ATAC-seq replicates and public 10x Genomics PBMC single-cell datasets, we demonstrate that *fp-tools* recovers biologically interpretable motif-centered aggregate patterns while preserving a reproducible analysis trail through versioned inputs, source tables, and regression-tested command outputs. *fp-tools* provides an integrated and auditable platform for ATAC-seq footprinting workflows and their motif-centered extensions.

Keywords: ATAC-seq; transcription factor footprinting; chromatin accessibility; motif analysis; pseudobulk ATAC-seq; replicate-aware reporting; reproducibility; command-line bioinformatics

1. Introduction

Chromatin accessibility profiling by the assay for transposase-accessible chromatin using sequencing (ATAC-seq) has become a standard readout of the active regulatory landscape of a cell [1]. Within accessible regions, sequence-specific transcription factors (TFs) protect the DNA they occupy from transposition. This creates local cut-site depletions, often called “footprints,” with higher accessibility in the nearby flanks. Detecting and interpreting these footprints supports inference of TF occupancy, differential regulation between conditions, and the functional consequences of non-coding variation. Earlier footprinting and occupancy models, including CENTIPEDE, Wellington, PIQ, and DeFCoM, established the value of combining sequence information with chromatin-accessibility profile shape for TFBS inference [2–5].

A mature but *fragmented* ecosystem of methods addresses pieces of this problem. ATAC-seq footprinting and differential-binding frameworks such as TOBIAS [6] and HINT-ATAC [7] correct enzymatic sequence bias and score footprints at a single TF-relevant scale. Multiscale and nucleosome-aware approaches, exemplified by the PRINT/scPrinter family [8], model accessibility structure across a range of footprint widths. Supervised TF-binding-site (TFBS) predictors such as maxATAC [9] and sequence-to-profile models

Received:

Accepted:

Published:

Copyright: © 2026 by the authors.

Submitted to *BioMedInformatics* for possible open access publication under the terms and conditions of the [Creative Commons Attribution \(CC BY\)](#) license.

such as ChromBPNet [10] learn occupancy from labeled data or raw sequence, and downstream interpretation tools and variant scorers translate models into motif and allele-level effects. Motif-discovery and motif-comparison tools, including MEME/STREME [11,12] and Tomtom [13], are also widely used after peak or footprint selection, but commonly run as separate steps with separate input and output conventions. Each tool typically ships its own interface, input conventions, and assumptions, so assembling a coherent, reproducible pipeline and rerunning analyses on common inputs remains laborious and error-prone.

Beyond this tooling friction, footprint interpretation has practical limitations that make reproducibility and careful signal handling important. ATAC-seq footprinting depends on base-resolution cut-site depletion and can be confounded by enzyme sequence bias, local accessibility, motif strength, and factor-specific binding kinetics [14]. Intrinsic cleavage bias remains a major confounder in accessibility-based footprinting, motivating explicit bias correction and careful interpretation [15]. Deep sequence-to-profile models and multiscale methods can improve sensitivity for some questions, but they often require substantial compute and specialized training workflows. For many applied ATAC-seq projects, the immediate need is therefore a transparent, repeatable workflow that keeps footprint interpretation traceable while producing biologically informative reports within one software ecosystem.

Three practical gaps motivate this work. First, **signal clarity**: users need to know which signal representation is appropriate for footprint calling and how normalization changes motif-centered interpretation. Second, **interface fragmentation**: moving from bias correction to footprint calling, motif matching, differential footprint analysis, and visualization commonly spans several packages with incompatible I/O. Third, **reproducible biological reporting**: many workflows produce intermediate tracks and tables but do not provide a consistent path to replicate-aware summaries, motif-centered aggregate views, de novo motif follow-up, and single-cell pseudobulk analyses.

We address these gaps with `fp-tools`, a reproducible ATAC-seq footprinting platform that connects bias correction, footprint calling, motif matching, differential footprint analysis, and aggregate visualization through a single command-line workflow. The same framework integrates de novo motif-discovery preparation, pseudobulk ATAC aggregation, replicate-aware comparison, quantile normalization, and interactive HTML reports. This paper describes the implemented methods, public-data validation examples, and reproducibility framework that tie these workflows together.

2. Materials and Methods

2.1. Package Architecture and Core Workflows

`fp-tools` implements performance-critical signal operations with Cython-backed routines. The workflow has four method layers: Tn5 sequence-bias correction, footprint scoring over corrected signal, motif-anchored differential TF-binding analysis, and aggregate visualization around TFBS or region sets. These layers share common region, signal, motif, and sample-label conventions so intermediate files can be passed through a complete analysis without format-specific glue code. YAML configuration files provide a reproducible way to save and rerun the same analysis without changing the underlying methods.

Following the signal model popularized by TOBIAS [6], `fp-tools` treats base-resolution, bias-corrected cut-site profiles as the primary footprinting signal. For footprint

calling, let x_i denote corrected Tn5 cut-site signal around candidate position p . The default footprint score is a center-versus-flank contrast,

$$F(p) = \frac{1}{|L| + |R|} \left(\sum_{i \in L} x_i + \sum_{i \in R} x_i \right) - \frac{1}{|C|} \sum_{i \in C} x_i, \quad (82)$$

where C is the protected center and L and R are flanking windows. Larger values therefore correspond to the expected footprint pattern: fewer cuts at the putative bound site and more cuts in nearby accessible flanks.

2.2. De Novo Motif Discovery Preparation

The de novo motif-discovery workflow begins from candidate footprint intervals, which may be produced by footprint calling or supplied by the user. For each candidate interval j with center p_j , strand a_j , and flank size h , the workflow exports a candidate-centered sequence

$$S_j = G[p_j - h, p_j + h], \quad (91)$$

reverse-complementing the sequence when strand information indicates the negative strand. The workflow records the input intervals, genome build, flank width, software versions, and MEME, DREME, STREME [11,12,16] or optional Tomtom [13] settings needed to reproduce each run. STREME was designed for accurate and versatile motif discovery from enriched sequence sets [12]. The same entry point can be used in two modes. In de novo-only mode, discovered motifs are reported as an independent motif set. In database-plus-de-novo mode, discovered motifs are matched to a known database such as JASPAR 2026 [17] and can be appended to the known motif set for downstream differential footprint analysis.

2.3. Pseudobulk Fragment Aggregation

The pseudobulk workflow groups single-cell ATAC fragments by a metadata column, such as cell type, cluster, donor, or treatment. Let F_b be the multiset of fragments assigned to barcode b , and let $g(b)$ be its group label. The pseudobulk fragment collection for group k is the concatenated multiset

$$F_k = \text{concat}_{b:g(b)=k}(F_b), \quad (106)$$

so repeated fragments and fragment multiplicity are preserved. Groups that pass user-defined cell-count and fragment-count filters are written as bulk-like fragment files with manifest records. When genome sizes are supplied, fp-tools can also emit CPM-normalized cut-site bigWigs,

$$y_{k,t} = 10^6 \frac{c_{k,t}}{\sum_u c_{k,u}}, \quad (111)$$

where $c_{k,t}$ is the cut-site count for group k at genomic position t . These tracks can be used directly by the same motif-centered aggregate plotting workflow used for bulk ATAC-seq data.

For visualization of cell-level marker footprint signatures, fp-tools also supports a K-nearest-neighbor (KNN) smoothing layer over single-cell fragments. For each marker TF, exact motif centers are scanned inside the selected accessible-bin universe. Fragment cut sites are counted in a ± 100 bp window around those motif centers for each barcode. Each cell is then assigned a local neighborhood profile by summing its 75 nearest neighbors in the single-cell embedding, and the same flank-minus-center footprint-signature score is computed from that neighborhood profile. Scores are oriented so that the expected marker cell type has the positive direction and then standardized per TF for UMAP display.

These KNN-smoothed signatures are used only as per-cell visualization scores; formal corrected footprint inference remains based on the grouped pseudobulk ATAC-correction and footprint-scoring workflow.

2.4. Replicate-aware Differential Footprint Analysis and Quantile Normalization

Replicate-aware analysis is integrated directly into the differential footprint comparison workflow. The differential comparison retains conventional motif-centered footprint-score logic while adding repeated condition labels, replicate-level summaries, and optional quantile normalization are added around the comparison. Repeated condition names, for example `-cond-names Bcell Bcell Tcell Tcell`, define biological replicate groups while preserving the original differential-footprint output structure. Let $x_{c,r,i}$ be the footprint score at site or background position i for replicate r of condition c .

With `-normalization sample-quantile`, each replicate is normalized before condition averaging. The normalization fit uses finite positive scores only, because negative or zero footprint-score values do not define a stable multiplicative quantile ratio. For sample s , `fp-tools` computes empirical score quantiles $q_s(p)$ on a 0.05–0.99 quantile grid and the reference quantile $q_*(p) = S^{-1} \sum_s q_s(p)$ across the S input samples. It then smooths the ratio $q_*(p)/q_s(p)$ and fits a sigmoid multiplicative factor $g_s(x)$; if the sample has too few positive values or the curve fit is ill-conditioned, a constant positive factor is used instead. During application, scores outside the fitted positive-score range use the nearest fitted range boundary for the normalization factor. Each score is transformed as

$$\tilde{x}_{s,i} = \max\{0, x_{s,i} g_s(x_{s,i})\}.$$

The final maximum operation clips transformed values at zero, preserving the interpretation that positive normalized values represent footprint-like center-versus-flank evidence. The `condition-quantile` mode first averages replicate background distributions within each condition, fits one factor $g_c(x)$ per condition, and applies that condition-level factor to all replicates in the condition. `none` disables quantile normalization. The same normalization helper is used by aggregate plotting, so differential comparison and motif-centered visualization are normalized consistently.

After normalization, condition-level scores are summarized as

$$\bar{x}_{c,i} = \frac{1}{n_c} \sum_{r=1}^{n_c} \tilde{x}_{c,r,i}, \quad s_{c,i} = \sqrt{\frac{1}{n_c - 1} \sum_{r=1}^{n_c} (\tilde{x}_{c,r,i} - \bar{x}_{c,i})^2}.$$

For a comparison between conditions c_1 and c_2 , `fp-tools` reports a footprint-score difference and a stabilized log fold change,

$$\Delta_i = \bar{x}_{c_1,i} - \bar{x}_{c_2,i}, \quad L_i = \log_2 \frac{\bar{x}_{c_1,i} + \epsilon}{\bar{x}_{c_2,i} + \epsilon},$$

where ϵ is estimated from the background signal unless supplied explicitly. When both conditions have replicate support, approximate standard errors are reported as descriptive summaries. Here, μ_1 and μ_2 denote the condition-level mean scores for the two conditions at the included motif sites, s_1 and s_2 denote the corresponding empirical SD estimates across replicates, and n_1 and n_2 are replicate counts:

$$\text{SE}(\Delta) = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}},$$

$$SE(L) \approx \frac{1}{\ln 2} \sqrt{\frac{s_1^2}{n_1(\mu_1 + \epsilon)^2} + \frac{s_2^2}{n_2(\mu_2 + \epsilon)^2}}.$$

For each motif and comparison, the reported change and pvalue columns are computed at the motif-reporting level. Candidate sites are included when either condition has a positive motif-site score, site-level log fold changes are averaged within peak intervals, and the observed motif statistic is defined as the difference between the fitted observed mean log fold change and the fitted background mean, scaled by the average of the observed and background SD estimates. `fp-tools` then simulates background change statistics from the fitted background distribution using the same number of observed intervals and applies a one-sample t-test comparing the simulated background changes with the observed motif statistic. P-values are adjusted across motifs within each contrast using Benjamini–Hochberg, yielding the `qvalue_bh` column and `significant_fdr05` flag. These p-values are screening statistics for the motif-centered report, especially at small replicate counts; the SE columns are reported separately as descriptive uncertainty summaries and are not themselves the hypothesis-test statistic. The `highlighted` column is retained only as a report-selection and visualization flag, not as a formal multiple-testing-corrected significance label. Additional replicate-oriented columns include replicate counts, condition score SDs, mean Δ , mean L , and standard-error summaries.

2.5. Public ATAC Processing and Reproducibility Provenance

The bulk replicate demonstration uses true biological replicates from the Buenrostro et al. ATAC-seq study [1]: GM12878/B-cell samples SRR891268 and SRR891269, and day-1 CD4 T-cell samples SRR891275 and SRR891276. Reads were trimmed with `fastp` [18], aligned with `Bowtie2` [19], processed with `SAMtools` [20] and `BEDTools` [21], and peaks were called with `MACS3` using the MACS model-based peak-calling framework [22]. Problematic genomic regions were filtered using the ENCODE blacklist [23]. Replicate peaks were merged into a common peak universe, bias-corrected, footprint-called, and analyzed by differential footprint comparison with and without sample-level quantile normalization. The analysis used JASPAR 2026 CORE vertebrates non-redundant motifs [17] for known-motif scanning. Exact command-line parameters and software versions are recorded in source TSV files in the manuscript `tables` directory and summarized in Tables S1 and S2.

2.6. Public Data, Demonstration Tiers, and Metrics

Public inputs are organized as versioned TSV manifests recording source, demonstration tier, cell type, donor, assay, accessions, assembly, output type, format, URL, checksum, status, local path, split, and notes. Discovery and download scripts query public repositories such as ENCODE [24] for released human GRCh38 ATAC-seq records, write resumable download reports, and keep raw public data out of version control. The tracked repository contains the scripts, manifests, schemas, small reference inputs, source tables, and figure code needed to rerun the manuscript analyses.

The demonstration design follows published recommendations for transparent and reproducible computational method evaluation [25,26], but it is not intended as an exhaustive head-to-head benchmark of all ATAC footprinting methods. Instead, it is organized as a tiered software-evaluation design (Table 1): a small-reference regression tier; a bulk ATAC replicate tier; a depth-robustness tier over downsampled ATAC; a de novo motif-preparation tier; and a real-data pseudobulk tier. Summary tables report command success, output integrity, runtime, memory, differential footprint summaries, motif-recovery summaries, and generated figures for the analyses included in this manuscript.

Table 1. Tiered benchmark design. Large data and outputs are untracked; only manifests, scripts, and compact reference results are committed.

Tier	Purpose	Primary outputs
Small reference	Installation and regression checks on compact inputs	command success, output existence, stable summaries
Bulk ATAC replicates	Public bulk ATAC replicate demonstrations for footprinting and differential reports	corrected tracks, footprint scores, motif reports, aggregate plots
Depth robustness	Shallow-depth sensitivity (50/25/10/5%)	downsampled workflow summaries and output checks
De novo motif prep	Candidate FASTA and motif-discovery validation	exported FASTA files and motif summaries
Pseudobulk	Single-cell extension and real-data demonstration	10x PBMC pseudobulk fragments, corrected footprint tracks, pairwise TF differential plots, KNN marker-signature UMAPs

2.7. Reproducibility Engineering208

The package ships a unit-test suite covering numerical helpers, CLI behavior, and deterministic golden CLI regressions for footprint calling, differential footprint analysis, and aggregate plotting, with an opt-in slow regression for bias correction. A continuous-integration workflow runs the suite on a clean checkout using compact reference inputs, builds source and wheel artifacts, and verifies the installed console scripts. The repository includes `environment.yml`, a Dockerfile, release metadata, and Make targets for test, build, and manuscript-smoke workflows; these choices follow FAIR-oriented research-software principles and established bioinformatics environment practices [27–29]. Detailed reviewer instructions are kept in `docs/reproduce-paper.md`. Benchmark scripts record random seeds, tool versions, command lines, and resolved manifests alongside the corresponding outputs. Figure source tables and editable figure files are retained for auditability.

3. Results220

We evaluate `fp-tools` as a software and reproducibility platform along three axes: workflow integration and command stability, controlled local engineering measurements, and biological plausibility in public bulk and single-cell ATAC-seq demonstrations.

3.1. A Unified, Reproducible Footprinting Platform224

`fp-tools` consolidates the core ATAC-seq footprinting steps into a single, consistent workflow with shared region/signal conventions, so a complete analysis, including bias correction, footprint calling, differential footprint analysis, and aggregate visualization, runs without crossing package boundaries or reformatting intermediates. YAML configuration records the same analysis in a machine-readable form for reruns and batch processing (Figure 1).

3.2. Interactive Interfaces and Report Review231

In addition to direct command-line execution, `fp-tools` provides browser-based interfaces for configuring analyses and reviewing results. The GUI writes reusable YAML configurations that remain compatible with command-line execution and provides guided panels for the same core commands shown in the command-line documentation (Figure 6). Differential-footprint and aggregate-batch workflows also produce self-contained HTML reports with embedded compressed data, searchable motif controls, editable group colors, synchronized aggregate profiles, and SVG export for publication-ready editing.




Figure 1. Overview of the fp-tools ATAC-seq footprinting platform. Bulk ATAC-seq BAM/fragments and single-cell ATAC fragments with annotations enter a shared command-oriented workflow. The workflow supports Tn5 bias correction, footprint calling, motif matching against curated databases, de novo motif discovery, replicate-aware differential footprint analysis, aggregate footprint visualization, and pseudobulk single-cell footprint analysis.

3.3. Regression-tested Workflow Outputs

The footprinting workflow is covered by deterministic regression tests that pin summary outputs for footprint calling, differential footprint analysis, and aggregate plotting. Public entry-point checks and YAML dry-runs verify that the command surface remains stable. These tests provide a reproducible contract for core analyses while allowing integrated modules to share the same input and output conventions.

3.4. *De Novo Motif Discovery Complements Database-based Motif Analysis*

The de novo motif-preparation path exports candidate-centered FASTA files, applies MEME, DREME, or STREME, optionally adds a Tomtom comparison against a known motif database, and summarizes motif results into TSV/HTML reports with inline consensus logos. This keeps motif discovery reproducible while relying on established motif-discovery algorithms. The workflow can be used alone to discover recurring patterns in candidate intervals, or after known-motif scanning to find additional motifs not represented in the selected database.

We therefore evaluated whether candidate-derived motifs add information beyond a fixed database using a 3-vs-3 public ENCODE K562 versus HepG2 ATAC-seq comparison [24]. Candidate footprints were called from bias-corrected cut-site tracks, condition-specific candidate FASTA files were exported with 75 bp flanks, and STREME was run reciprocally with K562 candidates against HepG2 controls and HepG2 candidates against K562 controls. These discovered motifs were then used in two modes. In de novo-only mode, the candidate-derived motifs were tested as their own motif set. In supplement mode, the same motifs were appended to a known motif database before running the sample-quantile-normalized differential-footprint analysis.

The full JASPAR 2026 vertebrate non-redundant set contained 1019 motifs, while the expanded reciprocal STREME runs contributed 500 candidate-derived motifs from both ENCODE cell lines. We therefore interpret this run as a site-recovery demonstration of how far de novo-only discovery can cover database-supported footprints, rather than as an exhaustive replacement for a curated motif database. In de novo-only mode, recovery was measured as base-pair overlap after merging the bound-site windows produced by the de novo scan and the corresponding JASPAR-supported scan. By this operational overlap metric, the discovered motifs recovered 84.0% of merged JASPAR-supported bound-site-window base pairs across the K562/HepG2 comparison, while 16.0% remained unique to the full JASPAR scan. The de novo-only scan also introduced additional candidate-derived bound-site windows, corresponding to 6.0% of the de novo-only merged window base pairs. In supplement mode, discovered motifs are appended to known motif sets before the same motif-aware differential footprint analysis. To create a sensitivity test, we also constructed a K562/HepG2-focused restricted JASPAR 2026 set with 867 motifs by removing common erythroid/K562 and liver/HepG2 motif families, then reran the comparison with and without the de novo motifs. In the interactive report, the restricted database yielded 92 display-priority motif families according to the highlighted report-selection flag; adding the 500 de novo motifs increased this browsing set to 153 families, including 24 de novo-derived motifs. Because highlighted is not a multiple-testing-corrected significance label, these counts are interpreted only as report-prioritization output, not as formal evidence for additional significant motif families. This sensitivity test illustrates the intended use case: when a chosen motif database is incomplete or conservative, candidate-derived motifs can be evaluated on their own and can also be carried forward into the same differential footprint report as database-supported motifs (Figure 2).

A representative K562 versus HepG2 differential footprint report is shown in Figure 3. The report combines selected motif logos, top differential motif rankings, a volcano summary, and replicate-level aggregate footprint profiles in one browser-review view.

3.5. *Replicate-aware Differential Footprint Analysis and Normalization*

The replicate-aware workflow summarizes differential footprint results with effect sizes, replicate support, and uncertainty estimates while keeping the comparison anchored to a common peak and motif-site universe. In the Buenrostro example, B-cell and T-cell




Figure 2. De novo motif discovery and validation workflow. (A) Candidate footprint or BED intervals are converted into centered FASTA sequences, analyzed with MEME/DREME/STREME, and summarized as motif tables and reports. (B) Merged site-window coverage compares shared and de novo-only motif-supported intervals in K562 and HepG2. (C,D) Representative K562 and HepG2 de novo motifs show footprint-like normalized cut-site aggregate profiles centered on motif instances.

replicates are compared after motif scanning on the merged peak set, so no additional peak alignment step is performed at the reporting stage.

Replicate-aware differential footprint analysis adds an interpretable layer to condition comparison: repeated condition names define biological replicate groups, condition means are reported alongside individual replicate profiles, and sample-level quantile normalization applies the same distribution-matching step before replicate group averages are compared. The no-normalization mode retains the baseline comparison without an additional sample-level distribution match, whereas the sample-quantile mode applies distribution matching before condition-level comparison and aggregate plotting.

As a biological consistency check for the same 2-vs-2 Buenrostro replicate comparison, we inspected aggregate corrected cut-site profiles for a small illustrative set selected to show visible center depletion and known or plausible immune-cell relevance. BACH2 showed a B/GM12878-biased pattern in this dataset while remaining biologically non-exclusive because it is also reused in T-cell states. Conversely, JUNB and ATF7 showed stronger




Figure 3. Interactive differential footprint report for K562 versus HepG2. The browser report summarizes selected differential motifs, sample styling controls, motif logos, top differential footprint-score bar plots, a volcano plot, and replicate-level motif-centered aggregate profiles. Blue indicates HepG2-biased footprint scores and red indicates K562-biased footprint scores.

T-cell aggregate depletion; JUNB is an AP-1 family regulator of T-cell differentiation [30], and ATF7 is included here as a strong T-cell-deeper aggregate example rather than a lineage-specific marker. We also include IRF4 because TCR signaling can drive graded IRF4 expression in T cells [31], although its aggregate depletion is modest in this dataset. We treat these panels as an illustrative sanity check rather than a marker-classification

benchmark; a more negative center-versus-flank profile indicates stronger footprint-like protection at motif centers.

3.6. Single-cell ATAC Fragments Support Pseudobulk Footprint Workflows

We evaluated pseudobulk aggregation using the public 10x Genomics PBMC granulocyte-depleted Multiome dataset [32]. The workflow used the official 10x ATAC fragments and GEX graph clusters, assigned 11,898 annotated cells to broad immune groups, and retained seven groups comprising 11,611 cells after requiring at least 300 cells and 50,000 fragments: B cells, CD4 T cells, NK/cytotoxic T cells, CD14 monocytes, FCGR3A monocytes, dendritic cells, and a mixed myeloid group. The prepared figure inputs used chromosomes chr1–chr22 plus chrX. The kept pseudobulks contained 13.6 to 144.0 million counted fragments per group and were written as indexed fragment files plus CPM-normalized cut-site bigWigs. The full pseudobulk-footprints route uses the same grouped fragments to emit pseudo-paired BAMs, run `atac-correct` with zero read shifting, score corrected footprint tracks, and optionally run motif-aware `diff-footprints` and aggregate plotting. Raw CPM cut-site tracks provide QC and chromatin-context signal alongside corrected footprint scores. These outputs were readable by the same aggregate-plotting code used for bulk ATAC data. This pseudobulk route is complementary to established single-cell chromatin-accessibility ecosystems such as ArchR and Signac and to motif-activity baselines such as chromVAR [33–35].

The pseudobulk footprint result table was also used to prioritize marker TFs for cell-level visualization. We summarized the three broad-cell pairwise comparisons as volcano plots and labeled only predeclared marker TFs from the two cell types included in each comparison when the differential-footprint direction matched the expected lineage direction. Third-lineage markers and opposite-direction markers were not annotated.

We also applied the per-cell KNN visualization layer to the public 10x Genomics 5k PBMC single-cell ATAC dataset [36]. The demonstration used the Cell Ranger ATAC 2.0.0 fragment dataset, broad B-cell, monocyte, and T/NK annotations derived from the single-cell embedding, and motif-centered marker scores for PAX5, CEBPB, TCF7, and four additional high-change marker TFs from the pairwise pseudobulk screen: CEBPA, SPIB, ZBTB7B, and POU2F2. For each cell, the plotted value is a KNN-smoothed footprint-signature score, not an independently bias-corrected single-cell footprint call. We summarize these scores as a marker-by-cell-type heatmap and as UMAP projections for one representative marker from each broad cell type. This visualization highlights expected marker directions while avoiding overinterpretation of sparse per-cell cut-site counts: PAX5 and POU2F2 show B-cell-associated signal, CEBPB and CEBPA show monocyte-associated signal, and TCF7 and ZBTB7B show T/NK-associated signal; SPIB is included as an additional high-change B-cell marker from the pseudobulk screen (Figure 4).

3.7. Benchmark and Figure Reproducibility

The tiered demonstration design (Table 1), manifest builders, schema validator, engineering-runtime recorder, and figure-generation scripts provide an executable path from public inputs to publication-ready outputs. These resources serve three purposes. First, they verify that the ATAC-seq footprinting workflow produces stable outputs on small fixtures and public-data subsets. Second, they support the manuscript demonstrations for bulk replicate processing, de novo motif preparation, pseudobulk aggregation, replicate-aware normalization, and report generation. Third, they keep public inputs, parameters, and generated figure sources auditable as the workflow is rerun after code changes.




Figure 4. Single-cell footprinting summary in public PBMC5k single-cell ATAC data. **(A)** KNN-smoothed per-cell footprint-signature heatmap for selected marker TFs and top cell-type-specific motif signatures across B cells, monocytes, and T/NK cells. **(B)** UMAP projections show broad cell-type annotations and representative marker footprint-signature z-scores. Positive values indicate the marker-oriented footprint-signature direction after per-TF standardization; formal corrected footprint inference is performed at the pseudobulk level.

The benchmark pipeline writes resolved manifests, software versions, parameters, and plotting tables next to generated figures. This makes the figures auditable: a reader can inspect which samples, motifs, regions, and signal tracks were used, and the same tier can be rerun after code changes. Representative local engineering measurements summarize runtime and peak memory for one local workflow comparison using the same input bundle and command-level outputs (Figure 5). The browser GUI emphasizes the same command-first workflow while making example loading, YAML review, run history,

and report inspection accessible to users who do not routinely work at the command line (Figure 6).




Figure 5. Representative local engineering comparison. (A) Runtime and (B) peak memory are summarized for one local workflow run comparing TOBIAS and fp-tools on the same analysis bundle. These measurements are included as an engineering check, not as an exhaustive cross-tool benchmark.




Figure 6. Browser GUI for command-compatible footprinting workflows. The GUI provides sidebar access to core commands, guided configuration panels, example loading, YAML preview, run-history review, and report/output inspection while preserving CLI-compatible configuration files.

4. Discussion

fp-tools contributes a reproducible platform for ATAC-seq footprinting and motif-centered regulatory analysis. Its comparative strengths are interface consistency across bias correction, footprint calling, differential footprint analysis, and aggregate visualization; deterministic regression coverage; explicit handling of replicate-aware summaries; and integrated support for de novo motif preparation and pseudobulk aggregation using the same region, motif, and signal conventions (Table 2). Together, these properties address the fragmentation that often separates footprint scoring, motif interpretation, and report generation.

Table 2. Qualitative scope comparison across widely used ATAC-seq regulatory-analysis tools. ✓, documented native support; ○, partial, indirect, or adjacent support; NA, outside the tool’s primary scope. The table summarizes workflow scope from public documentation and primary publications; it is not a performance benchmark and should be read with the different goals of these tools in mind.

Feature	fp-tools	TOBIAS	HINT-ATAC	PRINT/scPrinter	ChromBPNet	max ATAC
Bulk ATAC footprinting	✓	✓	✓	✓	○	NA
Tn5 bias correction	✓	✓	✓	✓	✓	NA
Footprint calling	✓	✓	✓	✓	NA	NA
Motif-aware differential analysis	✓	✓	○	○	NA	NA
De novo motif preparation	✓	NA	NA	✓	○	NA
scATAC/pseudobulk support	✓	○	○	✓	○	○
Replicate-aware reporting	✓	○	○	○	NA	NA
Visualization/reporting	✓	✓	○	✓	✓	○
YAML/batch execution	✓	NA	NA	NA	NA	NA

Deep sequence-to-profile models such as ChromBPNet provide powerful sequence-driven occupancy and variant-effect analyses [10]. Differential TF activity was also considered relative to orthogonal accessibility-based frameworks such as diffTF [37]. fp-tools addresses a complementary need: a lightweight, auditable ATAC-seq footprinting workflow that uses a conventional motif-centered footprinting model while integrating footprint calling, motif discovery, differential analysis, and reporting within one reproducible ecosystem. The replicate-aware summaries reported here use condition-level and replicate-level score variation to aid interpretation of differential footprint comparisons; they are designed for transparent reporting and visualization rather than as a replacement for specialized hierarchical models.

4.1. Threats to Validity and Scope

The present manuscript emphasizes reproducible workflow demonstrations, local engineering checks, and public-data examples; it is not a comprehensive cross-tool benchmark. Larger whole-genome transfer studies, including training on some cell lines and testing on others, and direct head-to-head runs of external tools such as TOBIAS, HINT-ATAC, PRINT/scPrinter, maxATAC, and ChromBPNet on identical locus sets will require larger controlled benchmark runs with repeated measurements, matched parameters, and explicit uncertainty. The runtime and memory measurements in Figure 5 are representative local engineering checks, not general performance claims.

Several biological panels are illustrative by design. The de novo recovery summary uses a merged-window overlap metric rather than an independent binding truth set, and the restricted-database sensitivity analysis depends on a specified motif-family removal list. The highlighted column is a report selection flag and should not be interpreted as statistical significance. The PBMC marker labels and aggregate examples emphasize expected-direction markers to show whether outputs are biologically plausible; full ranked tables are needed for formal marker-discovery claims. Finally, KNN-smoothed single-cell footprint-signature scores are visualization summaries over neighboring cells, whereas formal bias-corrected footprint inference is performed on grouped pseudobulk data.

Large public-benchmark inputs depend on external data acquisition and are therefore regenerated from committed manifests rather than stored directly in the repository. Natural next steps include larger external-tool comparisons, stronger per-replicate statistical models, additional public validation datasets, and expanded pseudobulk analyses across donors and tissues.

5. Conclusions

We presented `fp-tools`, a reproducible platform that unifies ATAC-seq footprinting with motif-discovery preparation, pseudobulk aggregation, and replicate-aware reporting. By pairing test-covered workflows with public-data manifests, source tables, and reproducible figure generation, the package supports standardized and auditable regulatory-genomics analysis. Future work will add larger whole-genome transfer studies, direct external-tool comparisons, multiscale and nucleosome-aware validation, and stronger hierarchical differential-binding validation.

Author Contributions: Conceptualization, Y.L. and C.Y.; methodology, Y.L.; software, Y.L.; validation, Y.L.; formal analysis, Y.L.; investigation, Y.L.; data curation, Y.L.; writing, original draft preparation, Y.L.; writing, review and editing, Y.L. and C.Y.; visualization, Y.L.; supervision, C.Y.; project administration, C.Y.; funding acquisition, C.Y. All authors have read and agreed to the published version of the manuscript.

Funding: This research received no external funding.

Institutional Review Board Statement: Not applicable. This study used only publicly available, de-identified datasets and did not involve new experiments on humans or animals.

Informed Consent Statement: Not applicable.

Data Availability Statement: `fp-tools` is openly available in the public GitHub repository [oncologylab/fp-tools](https://github.com/oncologylab/fp-tools); the source state for this manuscript revision is identified by the Git tag `manuscript-revision-2026-06-22`. The repository contains the benchmark manifests, schemas, analysis scripts, figure generators, source tables, and small result tables used in this manuscript. Public input data are obtained from the ENCODE Project [24], the 10x Genomics PBMC Multiome example [32], the 10x Genomics PBMC5k single-cell ATAC example [36], and motif databases such as JASPAR [17] using the included download and preparation scripts. Reviewer reproduction instructions are provided in `docs/reproduce-paper.md`; the computational environment is documented in `environment.yml` and the repository Dockerfile. Full raw public inputs are not stored in version control and are regenerated from committed manifests.

Acknowledgments: The authors thank the ENCODE Project Consortium and the maintainers of the open-source tools referenced in this work for making their data and software publicly available.

Conflicts of Interest: The authors declare no conflicts of interest.

Appendix

440

Table S1. Primary analysis parameters used for the public ATAC and pseudobulk workflows. The complete source table is provided in the Supplementary Materials as “Primary analysis parameters”.

Step	Parameters
Public input data	Bulk ATAC analyses use Buenrostro B-cell/T-cell FASTQs and ENCODE K562/HepG2 replicate tracks. Pseudobulk examples use public granulocyte-depleted 10x PBMC Multiome and PBMC5k scATAC fragment datasets. Motif analyses use JASPAR 2026 CORE vertebrate non-redundant motifs.
FASTQ preprocessing	Paired-end adapter detection and trimming with <code>fastp</code> ; HTML and JSON QC reports written for each sample.
Alignment and BAM processing	Reads aligned to hg38 with <code>bowtie2 -very-sensitive -X 2000</code> ; proper pairs retained with <code>samtools view -f 2 -F 2828 -q 30</code> ; BAMs sorted, fixed with <code>fixmate -m</code> , duplicate-marked with <code>markdup -r</code> , indexed, and summarized with <code>flagstat/idxstats</code> . Mitochondrial and nonstandard contigs excluded; hg38 blacklist v2 regions subtracted; MACS3 paired-end peaks called at q-value 0.01 and merged across replicates into one analysis peak set.
Blacklist and peak set	
Bias correction and footprint scores	<code>atac-correct</code> run on filtered BAMs with merged peaks and blacklist regions using default Tn5-bias settings; <code>call-footprints -score</code> footprint run on corrected cut-site bigWigs with default footprint and flank windows.
Differential footprint analysis	Buenrostro and ENCODE replicate footprint tracks analyzed with repeated condition labels, JASPAR 2026 motifs, and replicate reporting enabled. Runs include no-normalization and sample-quantile-normalized settings where indicated.
De novo motif validation	ENCODE K562/HepG2 candidate footprints exported as condition-level FASTA files with ± 75 bp flanks; reciprocal STREME 5.5.9 runs used <code>-dna -nmotifs 250</code> ; Tomtom 5.5.9 matched discovered motifs to JASPAR 2026; de novo motifs were evaluated alone and as additions to known motif sets.
Pseudobulk PBMC workflow	10x PBMC fragments grouped by broad immune labels, retaining groups with at least 300 cells and 50,000 fragments. The full <code>pseudobulk-footprints</code> route writes pseudo-BAMs, runs ATAC correction with zero read shifting for fragment-derived pseudobulk BAMs, and emits footprint-score tracks, motif-aware differential results, and aggregate plots. The PBMC5k demonstration uses broad B-cell/monocyte/T-NK labels, marker sites, pairwise volcano plots, and KNN-smoothed footprint-signature UMAPs.

Table S2. Software and resource versions used for manuscript analyses. The complete source table is provided in the Supplementary Materials as “Software and resource versions”.

Tool or resource	Version	Role
bedtools	2.31.1	Region sorting, merging, and blacklist subtraction
Biopython	1.87	Sequence and motif-related utilities
bowtie2	2.5.5	hg38 read alignment
Cell Ranger ARC	2.0.0	Source processing for the public 10x PBMC Multiome dataset
Cell Ranger ATAC	2.0.0	Source processing for the public 10x PBMC5k scATAC dataset
fastp	1.3.3	Paired-end FASTQ adapter detection, trimming, and QC
fp-tools	0.1.7	ATAC correction, footprint calling, motif matching, differential footprint analysis, pseudobulk utilities, and plotting helpers
hg38 blacklist	v2	Artifact-region filtering for processed bulk ATAC alignments
hg38 reference genome	GRCh38/hg38	Read alignment and sequence extraction
htslib	1.23.1	SAM/BAM/VCF I/O backend used by <code>samtools</code>
JASPAR CORE vertebrates	2026 non-redundant	Known motif database for differential footprint analysis
MACS3	3.0.4	ATAC-seq peak calling in paired-end mode
matplotlib	3.10.9	Manuscript figure rendering
MEME Suite, STREME, and Tomtom	5.5.9	De novo motif discovery and discovered-motif matching
NumPy	2.4.6	Numerical arrays and benchmark/figure calculations
pandas	2.3.3	Tabular processing for benchmark, manifest, and figure builders
pyBigWig	0.3.25	bigWig signal reading and writing
pysam	0.23.3	BAM and genomic interval I/O
Python	3.12.3	Runtime for analyses, figure builders, and provenance extraction
samtools	1.23.1	BAM filtering, sorting, fixmate, duplicate marking, indexing, and QC summaries
SciPy	1.17.1	Statistical and interpolation utilities

1. Buenrostro, J.D.; Giresi, P.G.; Zaba, L.C.; Chang, H.Y.; Greenleaf, W.J. Transposition of native chromatin for fast and sensitive epigenomic profiling of open chromatin, DNA-binding proteins and nucleosome position. *Nature Methods* **2013**, *10*, 1213–1218. <https://doi.org/10.1038/nmeth.2688>.
2. Pique-Regi, R.; Degner, J.F.; Pai, A.A.; Gaffney, D.J.; Gilad, Y.; Pritchard, J.K. Accurate inference of transcription factor binding from DNA sequence and chromatin accessibility data. *Genome Research* **2011**, *21*, 447–455. <https://doi.org/10.1101/gr.112623.110>.
3. Piper, J.; Elze, M.C.; Cauchy, P.; Cockerill, P.N.; Bonifer, C.; Ott, S. Wellington: a novel method for the accurate identification of digital genomic footprints from DNase-seq data. *Nucleic Acids Research* **2013**, *41*, e201. <https://doi.org/10.1093/nar/gkt850>.
4. Sherwood, R.I.; Hashimoto, T.; O'Donnell, C.W.; Lewis, S.; Barkal, A.A.; van Hoff, J.P.; Karun, V.; Jaakkola, T.; Gifford, D.K. Discovery of directional and nondirectional pioneer transcription factors by modeling DNase profile magnitude and shape. *Nature Biotechnology* **2014**, *32*, 171–178. <https://doi.org/10.1038/nbt.2798>.
5. Quach, B.; Furey, T.S. DeFCoM: analysis and modeling of transcription factor binding sites using a motif-centric genomic footprinter. *Bioinformatics* **2017**, *33*, 956–963. <https://doi.org/10.1093/bioinformatics/btw740>.
6. Bentsen, M.; Goymann, P.; Schultheis, H.; Klee, K.; Petrova, A.; Wiegandt, R.; Fust, A.; Preussner, J.; Kuenne, C.; Braun, T.; et al. ATAC-seq footprinting unravels kinetics of transcription factor binding during zygotic genome activation. *Nature Communications* **2020**, *11*, 4267. <https://doi.org/10.1038/s41467-020-18035-1>.
7. Li, Z.; Schulz, M.H.; Look, T.; Begemann, M.; Zenke, M.; Costa, I.G. Identification of transcription factor binding sites using ATAC-seq. *Genome Biology* **2019**, *20*, 45. <https://doi.org/10.1186/s13059-019-1642-2>.
8. Hu, Y.; Horlbeck, M.A.; Zhang, R.; Ma, S.; Shrestha, R.; Kartha, V.K.; Duarte, F.M.; Hock, C.; Savage, R.E.; Labade, A.; et al. Multiscale footprints reveal the organization of cis-regulatory elements. *Nature* **2025**, *638*, 779–786. <https://doi.org/10.1038/s41586-024-08443-4>.
9. Cazares, T.A.; Rizvi, F.W.; Iyer, B.; Chen, X.; Kotliar, M.; Bejjani, A.T.; Wayman, J.A.; Donmez, O.; Wronowski, B.; Parameswaran, S.; et al. maxATAC: Genome-scale transcription-factor binding prediction from ATAC-seq with deep neural networks. *PLoS Computational Biology* **2023**, *19*, e1010863. <https://doi.org/10.1371/journal.pcbi.1010863>.
10. Pampari, A.; Shcherbina, A.; Kvon, E.Z.; Kosicki, M.; Nair, S.; Kundu, S.; Kathiria, A.S.; Risca, V.I.; Hundenborn, K.; Zhou, Y.; et al. ChromBPNet: bias factorized, base-resolution deep learning models of chromatin accessibility reveal cis-regulatory sequence syntax, transcription factor footprints and regulatory variants. *bioRxiv* **2024**. Preprint, <https://doi.org/10.1101/2024.12.25.630221>.
11. Bailey, T.L.; Johnson, J.; Grant, C.E.; Noble, W.S. The MEME Suite. *Nucleic Acids Research* **2015**, *43*, W39–W49. <https://doi.org/10.1093/nar/gkv416>.
12. Bailey, T.L. STREME: accurate and versatile sequence motif discovery. *Bioinformatics* **2021**, *37*, 2834–2840. <https://doi.org/10.1093/bioinformatics/btab203>.
13. Gupta, S.; Stamatoyannopoulos, J.A.; Bailey, T.L.; Noble, W.S. Quantifying similarity between motifs. *Genome Biology* **2007**, *8*, R24. <https://doi.org/10.1186/gb-2007-8-2-r24>.
14. Sung, M.H.; Guertin, M.J.; Baek, S.; Hager, G.L. DNase footprint signatures are dictated by factor dynamics and DNA sequence. *Molecular Cell* **2014**, *56*, 275–285. <https://doi.org/10.1016/j.molcel.2014.08.016>.
15. Hu, S.S.; Liu, L.; Li, Q.; Ma, W.; Guertin, M.J. Intrinsic bias estimation for improved analysis of bulk and single-cell chromatin accessibility sequencing data using SELMA. *Nature Communications* **2022**, *13*, 6396. <https://doi.org/10.1038/s41467-022-33194-z>.
16. Bailey, T.L. DREME: motif discovery in transcription factor ChIP-seq data. *Bioinformatics* **2011**, *27*, 1653–1659. <https://doi.org/10.1093/bioinformatics/btr261>.
17. Ovek Baydar, D.; Rauluseviciute, I.; Aronsen, D.R.; Blanc-Mathieu, R.; Bonthuis, I.; de Beukelaer, H.; Ferenc, K.; Jegou, A.; Kumar, V.; Lemma, R.B.; et al. JASPAR 2026: expansion of transcription factor binding profiles and integration of deep learning models. *Nucleic Acids Research* **2026**, *54*, D184–D193. <https://doi.org/10.1093/nar/gkaf1209>.

18. Chen, S.; Zhou, Y.; Chen, Y.; Gu, J. fastp: an ultra-fast all-in-one FASTQ preprocessor. *Bioinformatics* **2018**, *34*, i884–i890. <https://doi.org/10.1093/bioinformatics/bty560>.
19. Langmead, B.; Salzberg, S.L. Fast gapped-read alignment with Bowtie 2. *Nature Methods* **2012**, *9*, 357–359. <https://doi.org/10.1038/nmeth.1923>.
20. Li, H.; Handsaker, B.; Wysoker, A.; Fennell, T.; Ruan, J.; Homer, N.; Marth, G.; Abecasis, G.; Durbin, R.; 1000 Genome Project Data Processing Subgroup. The Sequence Alignment/Map format and SAMtools. *Bioinformatics* **2009**, *25*, 2078–2079. <https://doi.org/10.1093/bioinformatics/btp352>.
21. Quinlan, A.R.; Hall, I.M. BEDTools: a flexible suite of utilities for comparing genomic features. *Bioinformatics* **2010**, *26*, 841–842. <https://doi.org/10.1093/bioinformatics/btq033>.
22. Zhang, Y.; Liu, T.; Meyer, C.A.; Eeckhoute, J.; Johnson, D.S.; Bernstein, B.E.; Nussbaum, C.; Myers, R.M.; Brown, M.; Li, W.; et al. Model-based Analysis of ChIP-Seq (MACS). *Genome Biology* **2008**, *9*, R137. <https://doi.org/10.1186/gb-2008-9-9-r137>.
23. Amemiya, H.M.; Kundaje, A.; Boyle, A.P. The ENCODE Blacklist: Identification of Problematic Regions of the Genome. *Scientific Reports* **2019**, *9*, 9354. <https://doi.org/10.1038/s41598-019-45839-z>.
24. The ENCODE Project Consortium. An integrated encyclopedia of DNA elements in the human genome. *Nature* **2012**, *489*, 57–74. <https://doi.org/10.1038/nature11247>.
25. Weber, L.M.; Hare, E.H.; Soneson, C.; Robinson, M.D. Essential guidelines for computational method benchmarking. *Genome Biology* **2019**, *20*, 125. <https://doi.org/10.1186/s13059-019-1738-8>.
26. Mangul, S.; Martin, L.S.; Hill, B.L.; Lam, A.K.M.; Distler, M.G.; Zelikovsky, A.; Eskin, E. Systematic benchmarking of omics computational tools. *Nature Communications* **2019**, *10*, 1393. <https://doi.org/10.1038/s41467-019-09406-4>.
27. Wilkinson, M.D.; Dumontier, M.; Aalbersberg, I.J.; et al. The FAIR Guiding Principles for scientific data management and stewardship. *Scientific Data* **2016**, *3*, 160018. <https://doi.org/10.1038/sdata.2016.18>.
28. Gruening, B.; Dale, R.; Sjoedin, A.; Chapman, B.A.; Rowe, J.; Tomkins-Tinch, C.H.; Valieris, R.; Koester, J. Bioconda: sustainable and comprehensive software distribution for the life sciences. *Nature Methods* **2018**, *15*, 475–476. <https://doi.org/10.1038/s41592-018-0046-7>.
29. Ewels, P.A.; Peltzer, A.; Fillinger, S.; Patel, H.; Alneberg, J.; Wilm, A.; Garcia, M.U.; Di Tommaso, P.; Nahnsen, S. The nf-core framework for community-curated bioinformatics pipelines. *Nature Biotechnology* **2020**, *38*, 276–278. <https://doi.org/10.1038/s41587-020-0439-x>.
30. Katagiri, T.; Kameda, H.; Nakano, H.; Yamazaki, S. Regulation of T cell differentiation by the AP-1 transcription factor JunB. *Immunological Medicine* **2021**, *44*, 197–203. <https://doi.org/10.1080/25785826.2021.1872838>.
31. Conley, J.M.; Gallagher, M.P.; Rao, A.; Berg, L.J. Activation of the Tec Kinase ITK Controls Graded IRF4 Expression in Response to Variations in TCR Signal Strength. *Journal of Immunology* **2020**, *205*, 335–345. <https://doi.org/10.4049/jimmunol.1900853>.
32. 10x Genomics. PBMC from a healthy donor, granulocytes removed through cell sorting, single cell multiome ATAC + gene expression. 10x Genomics public dataset, Cell Ranger ARC 2.0.0, 2021. ATAC fragments, peak calls, TF-analysis tables, and gene-expression clustering outputs downloaded from 10x Genomics sample files.
33. Granja, J.M.; Corces, M.R.; Pierce, S.E.; et al. ArchR is a scalable software package for integrative single-cell chromatin accessibility analysis. *Nature Genetics* **2021**, *53*, 403–411. <https://doi.org/10.1038/s41588-021-00790-6>.
34. Stuart, T.; Srivastava, A.; Madad, S.; Lareau, C.; Satija, R. Single-cell chromatin state analysis with Signac. *Nature Methods* **2021**, *18*, 1333–1341. <https://doi.org/10.1038/s41592-021-01282-5>.
35. Schep, A.N.; Wu, B.; Buenrostro, J.D.; Greenleaf, W.J. chromVAR: inferring transcription-factor-associated accessibility from single-cell epigenomic data. *Nature Methods* **2017**, *14*, 975–978. <https://doi.org/10.1038/nmeth.4401>.
36. 10x Genomics. 5k Peripheral Blood Mononuclear Cells (PBMCs) from a Healthy Donor (Next GEM v1.1). 10x Genomics public dataset, Cell Ranger ATAC 2.0.0, 2021. Single-cell ATAC fragments and Cell Ranger ATAC outputs for the atac_pbmc_5k_nextgem dataset.

37. Berest, I.; Arnold, C.; Reyes-Palomares, A.; Palla, G.; Rasmussen, K.D.; Giles, H.; Bruch, P.M.; Binder, H.; Guenther, T. Quantification of differential transcription factor activity and multiomics-based classification into activators and repressors: diffTF. *Cell Reports* **2019**, *29*, 3147–3159.e12. <https://doi.org/10.1016/j.celrep.2019.10.106>.